

Breast Cancer Diagnosis

Hongjie Liu, Xicheng Xie, Jiajun Tao, Shaohan Chen, Yujia Li

2023-02-27

Outline

- ▶ Background
- ▶ Task 1
- ▶ Task 2
- ▶ Task 3
- ▶ Task 4
- ▶ Discussions
- ▶ Limitations and Future Work
- ▶ Reference
- ▶ Q&A

Background

Breast Cancer Diagnosis

- ▶ The data is the breast cancer medical data retrieved from “breast-cancer.csv”, which has 569 rows and 32 columns.
- ▶ The first column ID labels individual breast tissue images; The second column Diagnosis identifies if the image is coming from cancer tissue or benign cases (M = malignant, B = benign). There are 357 benign and 212 malignant cases.
- ▶ The other 30 columns correspond to mean, standard deviation and the largest values (points on the tails) of the distributions of the following 10 features computed for the cellnuclei;

Objectives

- ▶ The goal of the exercise is to build a predictive model based on logistic regression to facilitate cancer diagnosis.
- ▶ We will move towards the goal with the steps of task 1, 2, 3 and 4.

Task 1 - Objective

Objective:

Build a logistic model to classify the images into malignant/benign, and write down your likelihood function, its gradient and Hessian matrix.

Task 1 - Build a logistic model

Define the “Diagnosis” variable will be coded as 1 for malignant cases and 0 for benign cases.

Given n i.i.d. observations with p predictors, we consider a logistic regression model

$$P(Y_i = 1 \mid \mathbf{x}_i) = \frac{e^{\mathbf{x}_i^\top \boldsymbol{\beta}}}{1 + e^{\mathbf{x}_i^\top \boldsymbol{\beta}}}, \quad i = 1, \dots, n \quad (1)$$

where $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^\top \in \mathbb{R}^{p+1}$ is the parameter vector, $\mathbf{x}_i = (1, X_{i1}, \dots, X_{ip})^\top$ is the vector of predictors in the i -th observation, and $Y_i \in \{0, 1\}$ is the binary response in the i -th observation.

Task 1 - Build a logistic model

Let $\mathbf{y} = (Y_1, Y_2, \dots, Y_n)^\top$ denote the response vector, and $\mathbf{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n)^\top \in \mathbb{R}^{n \times (p+1)}$ denote the design matrix.

The observed likelihood of $\{(Y_1, \mathbf{x}_1), (Y_2, \mathbf{x}_2), \dots, (Y_n, \mathbf{x}_n)\}$ is

$$L(\beta; \mathbf{y}, \mathbf{X}) = \prod_{i=1}^n \left[\left(\frac{e^{\mathbf{x}_i^\top \beta}}{1 + e^{\mathbf{x}_i^\top \beta}} \right)^{Y_i} \left(\frac{1}{1 + e^{\mathbf{x}_i^\top \beta}} \right)^{1-Y_i} \right]$$

Task 1 - Build a logistic model

Maximizing the likelihood is equivalent to maximizing the log-likelihood function:

$$f(\beta; \mathbf{y}, \mathbf{X}) = \sum_{i=1}^n \left[Y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right]. \quad (2)$$

The estimates of model parameters are

$$\hat{\beta} = \arg \max_{\beta} f(\beta; \mathbf{y}, \mathbf{X}),$$

and the optimization problem is

$$\max_{\beta} f(\beta; \mathbf{y}, \mathbf{X}). \quad (3)$$

Task 1 - Build a logistic model

Denote $p_i = P(Y_i = 1 \mid \mathbf{x}_i)$ as given in (1) and $\mathbf{p} = (p_1, p_2, \dots, p_n)^\top$. The gradient of f is:

$$\begin{aligned}\nabla f(\boldsymbol{\beta}; \mathbf{y}, \mathbf{X}) &= \mathbf{X}^\top (\mathbf{y} - \mathbf{p}) \\ &= \sum_{i=1}^n (Y_i - p_i) \mathbf{x}_i \\ &= \begin{pmatrix} \sum_{i=1}^n (Y_i - p_i) \\ \sum_{i=1}^n (Y_i - p_i) X_{i1} \\ \vdots \\ \sum_{i=1}^n (Y_i - p_i) X_{ip} \end{pmatrix}\end{aligned}$$

Task 1 - Build a logistic model

Denote $w_i = p_i(1 - p_i) \in (0, 1)$ and $\mathbf{W} = \text{diag}(w_1, \dots, w_n)$. The Hessian matrix of f is given by

$$\begin{aligned}\nabla^2 f(\beta; \mathbf{y}, \mathbf{X}) &= -\mathbf{X}^\top \mathbf{W} \mathbf{X} \\ &= -\sum_{i=1}^n w_i \mathbf{x}_i \mathbf{x}_i^\top \\ &= -\begin{pmatrix} \sum_{i=1}^n w_i & \sum_{i=1}^n w_i X_{i1} & \cdots & \sum_{i=1}^n w_i X_{i1} \\ \sum_{i=1}^n w_i X_{i1} & \sum_{i=1}^n w_i X_{i1}^2 & \cdots & \sum_{i=1}^n w_i X_{i1} X_{ip} \\ \vdots & \vdots & \ddots & \vdots \\ \sum_{i=1}^n w_i X_{ip} & \sum_{i=1}^n w_i X_{ip} X_{i1} & \cdots & \sum_{i=1}^n w_i X_{ip}^2 \end{pmatrix}\end{aligned}$$

Task 1 - Build a logistic model

Next, we show that the Hessian matrix $\nabla^2 f(\beta; \mathbf{y}, \mathbf{X})$ is a negative-definite matrix if \mathbf{X} has full rank.

Proof. For any $(p + 1)$ -dimensional nonzero vector α , given that \mathbf{X} has full rank, $\mathbf{X}\alpha$ is also a nonzero vector. Since \mathbf{W} is positive-definite, we have

$$\begin{aligned}\alpha^\top \nabla^2 f(\beta; \mathbf{y}, \mathbf{X}) \alpha &= \alpha^\top (-\mathbf{X}^\top \mathbf{W} \mathbf{X}) \alpha \\ &= -(\mathbf{X}\alpha)^\top \mathbf{W} (\mathbf{X}\alpha) \\ &< 0.\end{aligned}$$

Thus, $\nabla^2 f(\beta; \mathbf{y}, \mathbf{X})$ is negative-definite. □

Hence, the optimization problem (3) is a well-defined problem.

Task 1 - Logistic model R code

```
trainRows <- createDataPartition(y = cancer$diagnosis, p = 0.8, list = FALSE)
train <- cancer[trainRows, ]
test <- cancer[-trainRows, ]
glm.fit <- glm(diagnosis ~ .,
               data = train,
               subset = trainRows,
               family = binomial(link = "logit"))
summary(glm.fit)
pred <- predict(glm.fit, newdata = test, type = "response")
y_test <- factor(test$diagnosis)
auc_full <- auc(y_test, pred)
auc_full
```

Task 1 - Logistic model R code

Task 2 - Object

Objective:

Develop a Newton-Raphson algorithm to estimate your model

- Recall The target function f given in task 1 is:

$$f(\beta; \mathbf{y}, \mathbf{X}) = \sum_{i=1}^n \left[Y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right]. \quad (4)$$

- We develop a modified Newton-Raphson algorithm including a step-halving step to maximize the target function.

Task 2 - Newton-Raphson algorithm design

Algorithm 1 Newton-Raphson algorithm

Require: $f(\beta)$ - target function as given in (4); β_0 - starting value

Ensure: $\hat{\beta}$ such that $\hat{\beta} \approx \arg \max_{\beta} f(\beta)$

$i \leftarrow 0$, where i is the current number of iterations

$f(\beta_{-1}) \leftarrow -\infty$

while convergence criterion is not met **do**

$i \leftarrow i + 1$

$\mathbf{d}_i \leftarrow -[\nabla^2 f(\beta_{i-1})]^{-1} \nabla f(\beta_{i-1})$, where \mathbf{d}_i is the direction in the i -th iteration

$\lambda_i \leftarrow 1$, where λ_i is the multiplier in the i -th iteration

$\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$

while $f(\beta_i) \leq f(\beta_{i-1})$ **do**

$\lambda_i \leftarrow \lambda_i / 2$

$\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$

end while

end while

$\hat{\beta} \leftarrow \beta_i$

Task 2 - Newton-Raphson algorithm R code

The corresponding R-function for the algorithm

```
NewtonRaphson <- function(dat, func, start, tol = 1e-8, maxiter = 200) {  
  i <- 0  
  cur <- start  
  stuff <- func(dat, cur)  
  res <- c(0, stuff$f, cur)  
  prevf <- -Inf  
  X <- cbind(rep(1, nrow(dat)), as.matrix(dat[, -1]))  
  y <- dat[, 1]  
  warned <- 0  
  while (abs(stuff$f - prevf) > tol && i < maxiter) {  
    i <- i + 1  
    prevf <- stuff$f  
    prev <- cur  
    d <- -solve(stuff$Hess) %*% stuff$grad  
    cur <- prev + d  
    lambda <- 1  
    maxhalv <- 0  
    while (func(dat, cur)$f < prevf && maxhalv < 50) {  
      maxhalv <- maxhalv + 1  
      lambda <- lambda / 2  
      cur <- prev + lambda * d  
    }  
    stuff <- func(dat, cur)  
    res <- rbind(res, c(i, stuff$f, cur))  
    y_hat <- ifelse(X %*% cur > 0, 1, 0)  
    if (warned == 0 && sum(y - y_hat) == 0) {  
      warning("Complete separation occurs. Algorithm does not converge.")  
      warned <- 1  
    }  
  }  
  colnames(res) <- c("iter", "target_function", "(Intercept)", names(dat)[-1])  
  return(res)  
}
```


Task 2 - Complete separation

Sometimes our algorithm does not converge because of the complete separation.

A complete separation in a logistic regression, sometimes also referred as perfect prediction, occurs whenever there exists some vector of coefficients β such that $Y_i = 1$ whenever $x_i^\top \beta > 0$ and $Y_i = 0$ whenever $x_i^\top \beta \leq 0$.

Complete separation occurs when a linear function of predictors can perfectly classify the response.

Task 2 - Comparison

When there is no complete separation, the parameters output by `glm` function and our algorithm are the same.

Task 2 - Comparison

Comparison of using glm function and our algorithm (part1)

predictor	ours	glm
(Intercept)	111.7230206	90.9690365
radius_mean	-3646.8235387	-2560.3938902
texture_mean	-2.8949199	0.8812037
perimeter_mean	1257.9481173	789.2724398
area_mean	2091.5001210	1539.8345650
smoothness_mean	180.4591375	128.8762247
compactness_mean	-471.3749334	-346.6691873
concavity_mean	-0.3063034	9.0810082
concave.points_mean	287.3555092	215.4808031
symmetry_mean	10.8459784	10.4772590
fractal_dimension_mean	-40.6413497	-28.6916549
radius_se	854.4290933	617.5687358
texture_se	-105.8920782	-79.9788638
perimeter_se	-1231.6514585	-917.3916527
area_se	789.0550146	628.6222313

Task 2 - Comparison

Comparison of using glm function and our algorithm (part2)

predictor	ours	glm
smoothness_se	-83.34497	-63.76604
compactness_se	473.19705	344.31802
concavity_se	-440.66293	-323.85337
concave.points_se	503.37287	374.76109
symmetry_se	-144.45806	-108.62118
fractal_dimension_se	-462.42425	-339.16464
radius_worst	338.04088	197.93147
texture_worst	211.28178	155.07873
perimeter_worst	1406.29842	1068.10688
area_worst	-851.74931	-606.16579
smoothness_worst	-42.11738	-26.47595
compactness_worst	-480.79091	-346.40518
concavity_worst	518.45952	373.20892
concave.points_worst	-216.94150	-161.61771
symmetry_worst	98.05014	71.94885
fractal_dimension_worst	384.83305	282.82535

Task 3

Task 4

Discussions

- ▶ There is much freedom when designing the simulations.
- ▶ In our algorithm, we have 5 parameters. n , p , ratio, c , corr
- ▶ More parameters can be adjusted.

Limitations and Future Work

- ▶ Limitation: We reproduced high-dimensional scenarios, but we still don't know the solution.
- ▶ Future Work: We may adjust other parameters to investigate further.

Reference

1. Li Y, Hong HG, Ahmed SE, Li Y. Weak signals in high-dimensional regression: Detection, estimation and prediction. Appl Stochastic Models Bus Ind. 2018;1–16. <https://doi.org/10.1002/asmb.2340>

Q&A

- ▶ Thanks for listening!
- ▶ Any questions?