

Task 2

Task 2: Develop a Newton-Raphson algorithm to estimate your model.

The target function f given in task 1:

$$f(\beta; \mathbf{y}, \mathbf{X}) = \sum_{i=1}^n \left[Y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right]. \quad (1)$$

We develop a modified Newton-Raphson algorithm including a step-halving step. *(we probably don't need to ensure that the direction of the step is an ascent direction, since in this example Hessian is always negative-definite. but Hessian could be computationally singular when the starting points are bad)*

Algorithm 1 Newton-Raphson algorithm including a step-halving step

Require: $f(\beta)$ - target function as given in (1); β_0 - starting value

Ensure: $\hat{\beta}$ such that $\hat{\beta} \approx \arg \max_{\beta} f(\beta)$

```
1:  $i \leftarrow 0$ , where  $i$  is the current number of iterations
2:  $f(\beta_{-1}) \leftarrow -\infty$ 
3: while convergence criterion is not met do
4:    $i \leftarrow i + 1$ 
5:    $\mathbf{d}_i \leftarrow -[\nabla^2 f(\beta_{i-1})]^{-1} \nabla f(\beta_{i-1})$ , where  $\mathbf{d}_i$  is the direction in the  $i$ -th iteration
6:    $\lambda_i \leftarrow 1$ , where  $\lambda_i$  is the multiplier in the  $i$ -th iteration
7:    $\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$ 
8:   while  $f(\beta_i) \leq f(\beta_{i-1})$  do
9:      $\lambda_i \leftarrow \lambda_i / 2$ 
10:     $\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$ 
11:   end while
12: end while
13:  $\hat{\beta} \leftarrow \beta_i$ 
```

We write an **R**-function `NewtonRaphson` to implement the algorithm.

```
NewtonRaphson <- function(dat, func, start, tol = 1e-8, maxiter = 200) {
  i <- 0
  cur <- start
  stuff <- func(dat, cur)
  res <- c(0, stuff$f, cur)
  prevf <- -Inf
  X <- cbind(rep(1, nrow(dat)), as.matrix(dat[, -1]))
  y <- dat[, 1]
  warned <- 0
  while (abs(stuff$f - prevf) > tol && i < maxiter) {
    i <- i + 1
    prevf <- stuff$f
    prev <- cur
    d <- -solve(stuff$Hess) %*% stuff$grad
    cur <- prev + d
  }
}
```

```

lambda <- 1
maxhalv <- 0
while (func(dat, cur)$f < prevf && maxhalv < 50) {
  maxhalv <- maxhalv + 1
  lambda <- lambda / 2
  cur <- prev + lambda * d
}
stuff <- func(dat, cur)
res <- rbind(res, c(i, stuff$f, cur))
y_hat <- ifelse(X %>% cur > 0, 1, 0)
if (warned == 0 && sum(y - y_hat) == 0) {
  warning("Complete separation occurs. Algorithm does not converge.")
  warned <- 1
}
}
colnames(res) <- c("iter", "target_function", "(Intercept)", names(dat)[-1])
return(res)
}

```

Data preprocessing and data partition.

```

bc_df <- read.csv("breast-cancer.csv")[-c(1, 33)] %>% # remove variable ID and an NA column
  mutate(diagnosis = ifelse(diagnosis == "M", 1, 0)) # code malignant cases as 1
bc_df[, -1] <- scale(bc_df[, -1]) # predictors are standardized for the logistic-LASSO model in task 3

set.seed(1)
indexTrain <- createDataPartition(y = bc_df$diagnosis, p = 0.8, list = FALSE)
Training <- bc_df[indexTrain, ]
Test <- bc_df[-indexTrain, ]

glm(diagnosis ~ ., family = binomial(link = "logit"), data = Training)

```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

##
## Call:  glm(formula = diagnosis ~ ., family = binomial(link = "logit"),
##       data = Training)
##
## Coefficients:
##           (Intercept)           radius_mean           texture_mean
##             90.9690             -2560.3939              0.8812
##      perimeter_mean           area_mean           smoothness_mean
##       789.2724             1539.8346             128.8762
## compactness_mean           concavity_mean concave.points_mean
##      -346.6692              9.0810             215.4808
##      symmetry_mean fractal_dimension_mean           radius_se
##       10.4773             -28.6917             617.5687
##           texture_se           perimeter_se           area_se
##      -79.9789             -917.3917             628.6222
##      smoothness_se           compactness_se           concavity_se

```

```
##           -63.7660           344.3180           -323.8534
##      concave.points_se      symmetry_se      fractal_dimension_se
##           374.7611           -108.6212           -339.1646
##           radius_worst      texture_worst      perimeter_worst
##           197.9315           155.0787           1068.1069
##           area_worst      smoothness_worst      compactness_worst
##           -606.1658           -26.4759           -346.4052
##           concavity_worst      concave.points_worst      symmetry_worst
##           373.2089           -161.6177           71.9489
## fractal_dimension_worst
##           282.8254
##
## Degrees of Freedom: 455 Total (i.e. Null);  425 Residual
## Null Deviance:      601.3
## Residual Deviance: 1.311e-06      AIC: 62
```

```
logisticstuff <- function(dat, betavec) {
  dat <- as.matrix(dat)
  n <- nrow(dat)
  p <- ncol(dat) - 1
  X <- cbind(rep(1, n), dat[, -1]) # design matrix
  y <- dat[, 1] # response vector
  u <- X %*% betavec #  $x_i^T \beta$ ,  $i=1, \dots, n$ 
  f <- sum(y * u - log1pexp(u)) # function `log1pexp` to compute  $\log(1 + \exp(x))$ 
  p_vec <- sigmoid(u) # function `sigmoid` to compute  $\exp(x)/(1 + \exp(x))$ 
  grad <- t(X) %*% (y - p_vec)
  Hess <- -t(X) %*% diag(c(p_vec * (1 - p_vec))) %*% X
  return(list(f = f, grad = grad, Hess = Hess))
}
```

We fit a logistic regression model on the training data using our NewtonRaphson function.

```
res <- NewtonRaphson(dat = Training, func = logisticstuff, start = rep(0, ncol(Training)))

## Warning in NewtonRaphson(dat = Training, func = logisticstuff, start = rep(0, :
## Complete separation occurs. Algorithm does not converge.
```

```
tail(res)
```

```
## iter target_function (Intercept) radius_mean texture_mean perimeter_mean
## 26 -3.235158e-07 89.24826 -2822.544 -1.373617 960.4715
## 27 -1.190284e-07 94.06615 -2977.405 -1.518686 1010.5839
## 28 -4.379281e-08 98.80114 -3134.863 -1.704236 1063.1357
## 29 -1.611204e-08 103.40055 -3296.525 -1.956417 1119.6337
## 30 -5.927773e-09 107.76671 -3465.377 -2.323911 1182.8800
## 31 -2.180833e-09 111.72302 -3646.824 -2.894920 1257.9481
## area_mean smoothness_mean compactness_mean concavity_mean concave.points_mean
## 1628.129 140.0014 -368.0153 2.2563839 224.3770
## 1719.801 147.7362 -388.2773 2.1569756 236.9795
## 1811.729 155.5634 -408.6480 1.9513295 249.5825
## 1904.079 163.5407 -429.1918 1.5722084 262.1834
## 1997.154 171.7748 -450.0285 0.8948694 274.7774
```

```
##      2091.500      180.4591      -471.3749      -0.3063034      287.3555
## symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
##      8.431909      -31.04890      664.0699      -82.56232      -964.1721
##      9.125769      -32.85952      700.8645      -87.29044      -1017.2613
##      9.763178      -34.69764      737.9928      -92.00209      -1070.4606
##      10.311347      -36.58090      775.6589      -96.68750      -1123.8325
##      10.709177      -38.54224      814.2424      -101.32850      -1177.4929
##      10.845978      -40.64135      854.4291      -105.89208      -1231.6515
## area_se smoothness_se compactness_se concavity_se concave.points_se
## 627.0977      -65.29781      368.1306      -343.4408      393.0646
## 661.3045      -69.00116      388.5425      -362.5350      414.9751
## 695.0465      -72.68019      409.1130      -381.7210      436.9259
## 728.0267      -76.31976      429.9373      -401.0546      458.9405
## 759.6946      -79.89169      451.1922      -420.6395      481.0630
## 789.0550      -83.34497      473.1970      -440.6629      503.3729
## symmetry_se fractal_dimension_se radius_worst texture_worst perimeter_worst
## -113.2103      -360.5082      247.8024      164.0343      1108.082
## -119.4203      -380.3806      263.6037      173.3317      1168.164
## -125.6423      -400.3884      279.9778      182.6679      1228.146
## -131.8815      -420.6120      297.2940      192.0677      1287.947
## -138.1475      -441.2006      316.2365      201.5777      1347.420
## -144.4581      -462.4243      338.0409      211.2818      1406.298
## area_worst smoothness_worst compactness_worst concavity_worst
## -659.5481      -32.14856      -373.8623      401.9769
## -696.8089      -33.84766      -394.4311      424.5118
## -734.3769      -35.62248      -415.2146      447.2448
## -772.4418      -37.52094      -436.3424      470.2984
## -811.3548      -39.63214      -458.0558      493.8996
## -851.7493      -42.11738      -480.7909      518.4595
## concave.points_worst symmetry_worst fractal_dimension_worst
## -169.3830      77.46596      300.1686
## -178.8370      81.34102      316.7304
## -188.3071      85.28220      333.3894
## -197.8017      89.32603      350.2048
## -207.3369      93.54077      367.2864
## -216.9415      98.05014      384.8331
```

Our function also does not converge, because a complete separation occurs. A complete separation in a logistic regression, sometimes also referred as perfect prediction, which occurs whenever there exists some vector of coefficients β such that $Y_i = 1$ whenever $\mathbf{x}_i^\top \beta > 0$ and $Y_i = 0$ whenever $\mathbf{x}_i^\top \beta \leq 0$. In other words, complete separation occurs whenever a linear function of predictors can generate perfect predictions of response.

```
X <- cbind(rep(1, nrow(Training)), model.matrix(diagnosis ~ ., Training)[, -1])
y <- Training$diagnosis
coef_newton <- res[nrow(res), -c(1, 2)]
y_hat <- ifelse(X %*% coef_newton > 0, 1, 0) # predictions
sum(y - y_hat) # complete separation
```

```
## [1] 0
```

We can prove that: when there exists a vector of coefficients $\hat{\beta}$ such that $Y_i = 1$ whenever $\mathbf{x}_i^\top \hat{\beta} > 0$ and $Y_i = 0$ whenever $\mathbf{x}_i^\top \hat{\beta} \leq 0$, there does not exist $\beta^* \in \mathbb{R}^{(p+1)}$ such that $\beta^* = \arg \max_{\beta} f(\beta)$, where f is given in (1). Thus our algorithm does not converge.

Proof.

Assume such β^* exists, then $\forall \beta \in \mathbb{R}^{p+1}$, we have $f(\beta) \leq f(\beta^*)$.

First, we prove that: there exists a vector of coefficients $\tilde{\beta}$ such that $Y_i = 1$ whenever $\mathbf{x}_i^\top \tilde{\beta} > 0$ and $Y_i = 0$ whenever $\mathbf{x}_i^\top \tilde{\beta} < 0$.

Let $A_1 = \{i : Y_i = 1\} = \{i : \mathbf{x}_i^\top \hat{\beta} > 0\}$ and $A_0 = \{i : Y_i = 0\} = \{i : \mathbf{x}_i^\top \hat{\beta} \leq 0\}$. Then we have

$$\epsilon := \min_{i \in A_1} (\mathbf{x}_i^\top \hat{\beta}) > 0.$$

Let $\tilde{\beta} = \hat{\beta} - (\epsilon/2, 0, \dots, 0)^\top$. Given that $X_{i1} = 1$ for all i , we have

$$\begin{aligned} \mathbf{x}_i^\top \tilde{\beta} &= \mathbf{x}_i^\top \hat{\beta} - \epsilon/2 \cdot 1 \geq \epsilon - \epsilon/2 = \epsilon/2 > 0, \quad \forall i \in A_1 \\ \mathbf{x}_i^\top \tilde{\beta} &= \mathbf{x}_i^\top \hat{\beta} - \epsilon/2 \cdot 1 \leq 0 - \epsilon/2 = -\epsilon/2 < 0, \quad \forall i \in A_0 \end{aligned}$$

Thus we have $A_1 = \{i : Y_i = 1\} = \{i : \mathbf{x}_i^\top \tilde{\beta} > 0\}$ and $A_0 = \{i : Y_i = 0\} = \{i : \mathbf{x}_i^\top \tilde{\beta} < 0\}$.

Next, we prove that

$$\lim_{k \rightarrow \infty} f(k\tilde{\beta}) = 0. \quad (2)$$

$\forall k > 0$, we have $A_1 = \{i : \mathbf{x}_i^\top (k\tilde{\beta}) > 0\}$ and $A_0 = \{i : \mathbf{x}_i^\top (k\tilde{\beta}) < 0\}$.

Thus,

$$\begin{aligned} \lim_{k \rightarrow \infty} f(k\tilde{\beta}) &= \lim_{k \rightarrow \infty} \sum_{i \in A_1} \left[Y_i \mathbf{x}_i^\top (k\tilde{\beta}) - \log \left(1 + e^{\mathbf{x}_i^\top (k\tilde{\beta})} \right) \right] + \lim_{k \rightarrow \infty} \sum_{i \in A_0} \left[Y_i \mathbf{x}_i^\top (k\tilde{\beta}) - \log \left(1 + e^{\mathbf{x}_i^\top (k\tilde{\beta})} \right) \right] \\ &= \sum_{i \in A_1} \lim_{k \rightarrow \infty} \left[k \mathbf{x}_i^\top \tilde{\beta} - \log \left(1 + e^{k \mathbf{x}_i^\top \tilde{\beta}} \right) \right] + \sum_{i \in A_0} \lim_{k \rightarrow \infty} \left[-\log \left(1 + e^{k \mathbf{x}_i^\top \tilde{\beta}} \right) \right] \\ &= \sum_{i \in A_1} \lim_{z \rightarrow \infty} [z - \log(1 + e^z)] + \sum_{i \in A_0} (-\log 1) \\ &= 0 + 0 = 0. \end{aligned}$$

Last, we prove that: there exists $\beta \in \mathbb{R}^{p+1}$ such that $f(\beta) > f(\beta^*)$, which is contradictory to the statement that $\forall \beta \in \mathbb{R}^{p+1}$, $f(\beta) \leq f(\beta^*)$.

Note that $f(\beta) < 0$ holds for any $\beta \in \mathbb{R}$, then we have $f(\beta^*) < 0$.

Given that $f(\beta^*) < 0$ and (2) holds, there exists $N \in \mathbb{R}$ such that $\forall k > N$, $f(k\tilde{\beta}) > f(\beta^*)$.

Thus our assumption must be false. □

We compare the results of using the `glm` function and our `NewtonRaphson` function. (meaningless, since both do not converge)

```
tibble(
  predictor = c("(Intercept)", names(Training)[-1]),
  ours = res[nrow(res), -c(1, 2)],
  glm = glm(diagnosis ~ ., family = binomial(link = "logit"), data = Training)$coefficients
) %>%
  knitr::kable()
```

predictor	ours	glm
(Intercept)	111.7230205	90.9690365
radius_mean	-3646.8235396	-2560.3938902

predictor	ours	glm
texture_mean	-2.8949199	0.8812037
perimeter_mean	1257.9481182	789.2724398
area_mean	2091.5001211	1539.8345650
smoothness_mean	180.4591376	128.8762247
compactness_mean	-471.3749334	-346.6691873
concavity_mean	-0.3063034	9.0810082
concave.points_mean	287.3555092	215.4808031
symmetry_mean	10.8459784	10.4772590
fractal_dimension_mean	-40.6413497	-28.6916549
radius_se	854.4290934	617.5687358
texture_se	-105.8920782	-79.9788638
perimeter_se	-1231.6514585	-917.3916527
area_se	789.0550144	628.6222313
smoothness_se	-83.3449730	-63.7660367
compactness_se	473.1970464	344.3180183
concavity_se	-440.6629326	-323.8533730
concave.points_se	503.3728692	374.7610875
symmetry_se	-144.4580558	-108.6211792
fractal_dimension_se	-462.4242540	-339.1646360
radius_worst	338.0408767	197.9314738
texture_worst	211.2817780	155.0787339
perimeter_worst	1406.2984247	1068.1068808
area_worst	-851.7493135	-606.1657870
smoothness_worst	-42.1173805	-26.4759499
compactness_worst	-480.7909132	-346.4051824
concavity_worst	518.4595221	373.2089180
concave.points_worst	-216.9414972	-161.6177074
symmetry_worst	98.0501371	71.9488540
fractal_dimension_worst	384.8330525	282.8253512

We probably won't conduct resampling. Ignore the below.

Resampling on training data: *Does the following resampling method work?*

```
?caret::resamples
```

Hothorn et al. The design and analysis of benchmark experiments. Journal of Computational and Graphical Statistics (2005) vol. 14 (3) pp. 675-699

<https://ro.uow.edu.au/cgi/viewcontent.cgi?article=3494&context=commpapers>

RW-OOB

```
B = 100 # number of bootstrap samples
set.seed(1)
auc.logit <- rep(NA, B)
for (i in 1:B) {
  index_bs <- sample(nrow(Training), replace = TRUE)
  sample <- Training[index_bs, ]
  out <- Training[-index_bs, ]
  res <- NewtonRaphson(dat = sample, func = logisticstuff, start = rep(0, ncol(sample)))
  betavec <- res[nrow(res), 3:ncol(res)]
  X <- cbind(rep(1, nrow(out)), model.matrix(diagnosis ~ ., out)[, -1])
```

```

y <- out$diagnosis
u <- X %*% betavec
phat <- sigmoid(u)[, 1]
roc <- roc(response = y, predictor = phat)
auc <- roc$auc[1]
auc.logit[i] <- auc
}
summary(auc.logit)

```

```

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.9117 0.9518 0.9622 0.9611 0.9703 0.9933

```

```
boxplot(auc.logit)
```

