

Task 2

Task 2: Develop a Newton-Raphson algorithm to estimate your model.

The target function f given in task 1:

$$f(\beta; \mathbf{y}, \mathbf{X}) = \sum_{i=1}^n \left[Y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right]. \quad (1)$$

We develop a modified Newton-Raphson algorithm including a step-halving step. *(we probably don't need to ensure that the direction of the step is an ascent direction, since in this example Hessian is always negative-definite. but Hessian could be computationally singular when the starting points are bad)*

Algorithm 1 Newton-Raphson algorithm including a step-halving step

Require: $f(\beta)$ - target function as given in (1); β_0 - starting value

Ensure: $\hat{\beta}$ such that $\hat{\beta} \approx \arg \max_{\beta} f(\beta)$

```
1:  $i \leftarrow 0$ , where  $i$  is the current number of iterations
2:  $f(\beta_{-1}) \leftarrow -\infty$ 
3: while convergence criterion is not met do
4:    $i \leftarrow i + 1$ 
5:    $\mathbf{d}_i \leftarrow -[\nabla^2 f(\beta_{i-1})]^{-1} \nabla f(\beta_{i-1})$ , where  $\mathbf{d}_i$  is the direction in the  $i$ -th iteration
6:    $\lambda_i \leftarrow 1$ , where  $\lambda_i$  is the multiplier in the  $i$ -th iteration
7:    $\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$ 
8:   while  $f(\beta_i) \leq f(\beta_{i-1})$  do
9:      $\lambda_i \leftarrow \lambda_i / 2$ 
10:     $\beta_i \leftarrow \beta_{i-1} + \lambda_i \mathbf{d}_i$ 
11:   end while
12: end while
13:  $\hat{\beta} \leftarrow \beta_i$ 
```

We write an **R**-function `NewtonRaphson` to implement the algorithm.

```
NewtonRaphson <- function(dat, func, start, tol = 1e-10) {
  i <- 0
  cur <- start
  stuff <- func(dat, cur)
  res <- c(0, stuff$f, cur)
  prevf <- -Inf
  while (abs(stuff$f - prevf) > tol) {
    i <- i + 1
    prevf <- stuff$f
    prev <- cur
    d <- -solve(stuff$Hess) %% stuff$grad
    cur <- prev + d
    lambda <- 1
    maxhalv <- 0
    while (func(dat, cur)$f < prevf && maxhalv < 50) {
```

```

    maxhalv <- maxhalv + 1
    lambda <- lambda / 2
    cur <- prev + lambda * d
  }
  stuff <- func(dat, cur)
  res <- rbind(res, c(i, stuff$f, cur))
}
colnames(res) <- c("iter", "target_function", "(Intercept)", names(dat)[-1])
return(res)
}

```

Data preprocessing and data partition.

```

bc_df <- read.csv("breast-cancer.csv")[-c(1, 33)] %>% # remove variable ID and an NA column
  mutate(diagnosis = ifelse(diagnosis == "M", 1, 0)) # code malignant cases as 1
bc_df[, -1] <- scale(bc_df[, -1]) # predictors are standardized for the logistic-LASSO model in task 3

set.seed(1)
indexTrain <- createDataPartition(y = bc_df$diagnosis, p = 0.8, list = FALSE)
Training <- bc_df[indexTrain, ]
Test <- bc_df[-indexTrain, ]

# correlation coefficients close to 1. need to remove some variables to make the algorithm converge
glm(diagnosis ~ ., family = binomial(link = "logit"), data = Training)

```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

##
## Call:  glm(formula = diagnosis ~ ., family = binomial(link = "logit"),
##       data = Training)
##
## Coefficients:
##           (Intercept)           radius_mean           texture_mean
##           90.9690           -2560.3939           0.8812
##           perimeter_mean           area_mean           smoothness_mean
##           789.2724           1539.8346           128.8762
##           compactness_mean           concavity_mean           concave.points_mean
##           -346.6692           9.0810           215.4808
##           symmetry_mean           fractal_dimension_mean           radius_se
##           10.4773           -28.6917           617.5687
##           texture_se           perimeter_se           area_se
##           -79.9789           -917.3917           628.6222
##           smoothness_se           compactness_se           concavity_se
##           -63.7660           344.3180           -323.8534
##           concave.points_se           symmetry_se           fractal_dimension_se
##           374.7611           -108.6212           -339.1646
##           radius_worst           texture_worst           perimeter_worst
##           197.9315           155.0787           1068.1069
##           area_worst           smoothness_worst           compactness_worst
##           -606.1658           -26.4759           -346.4052

```

```
##          concavity_worst      concave.points_worst      symmetry_worst
##                373.2089                -161.6177                71.9489
## fractal_dimension_worst
##                282.8254
##
## Degrees of Freedom: 455 Total (i.e. Null); 425 Residual
## Null Deviance:      601.3
## Residual Deviance: 1.311e-06      AIC: 62
```

Remove some variables. Here we select all mean predictors as an example (still have highly-correlated variables). **(Should remove variables that are highly correlated. NOT DECIDED!)**

```
# select some variables. should decide which variables to choose
useful <- names(bc_df)[1:11] # e.g., select all mean variables
bc_df2 <-
  bc_df %>%
  select(all_of(useful))

set.seed(1)
indexTrain <- createDataPartition(y = bc_df2$diagnosis, p = 0.8, list = FALSE)
Training <- bc_df2[indexTrain, ]
Test <- bc_df2[-indexTrain, ]

glm(diagnosis ~ ., family = binomial(link = "logit"), data = Training)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
##
## Call: glm(formula = diagnosis ~ ., family = binomial(link = "logit"),
##      data = Training)
##
## Coefficients:
##      (Intercept)      radius_mean      texture_mean
##           0.4117        -18.3611           1.6021
##    perimeter_mean      area_mean      smoothness_mean
##           8.6626          15.1088           1.0744
##    compactness_mean    concavity_mean    concave.points_mean
##          -0.0888           0.3343           2.2067
##      symmetry_mean    fractal_dimension_mean
##           0.6342          -0.5820
##
## Degrees of Freedom: 455 Total (i.e. Null); 445 Residual
## Null Deviance:      601.3
## Residual Deviance: 116.7      AIC: 138.7
```

```
logisticstuff <- function(dat, betavec) {
  dat <- as.matrix(dat)
  n <- nrow(dat)
  p <- ncol(dat) - 1
  X <- cbind(rep(1, n), dat[, -1]) # design matrix
  y <- dat[, 1] # response vector
  u <- X %*% betavec # x_i^T beta, i=1,...,n
  f <- sum(y * u - log1pexp(u)) # function `log1pexp` to compute log(1 + exp(x))
```

```

p_vec <- sigmoid(u) # function `sigmoid` to compute  $\exp(x)/(1 + \exp(x))$ 
grad <- t(X) %*% (y - p_vec)
Hess <- -t(X) %*% diag(c(p_vec * (1 - p_vec))) %*% X
return(list(f = f, grad = grad, Hess = Hess))
}

```

We fit a logistic regression model on the training data using our `NewtonRaphson` function.

```

res <- NewtonRaphson(dat = Training, func = logisticstuff, start = rep(0, ncol(Training)))
res

```

```

##      iter target_function (Intercept) radius_mean texture_mean perimeter_mean
## res    0      -316.07511    0.0000000    0.0000000    0.0000000    0.0000000
##      1      -129.88900   -0.5213358    7.46247726    0.3593164    -6.2762894
##      2       -86.99235   -0.7300300   10.08483576    0.6391330    -8.2532785
##      3       -69.28577   -0.8159302    7.59038364    0.9534132    -5.9835656
##      4       -62.07699   -0.6613894   -0.09853138    1.2435119   -0.2505531
##      5       -59.17237   -0.1869629   -9.84136501    1.4404154    5.1385186
##      6       -58.39991    0.2759214  -16.39521429    1.5559376    7.8043536
##      7       -58.35396    0.4048031  -18.24786635    1.5989307    8.6072512
##      8       -58.35380    0.4116856  -18.36070646    1.6021338    8.6623617
##      9       -58.35380    0.4117085  -18.36110260    1.6021463    8.6625595
##     10       -58.35380    0.4117085  -18.36110260    1.6021463    8.6625595
##      area_mean smoothness_mean compactness_mean concavity_mean
## res 0.0000000    0.0000000    0.0000000    0.0000000
##     -1.0612656    0.1067701    0.19577818    0.1839443
##     -1.4030153    0.2615841    0.28540534    0.3765710
##     -0.4711325    0.4832265    0.21409284    0.5105891
##      2.6852622    0.7505353    0.04538621    0.5466722
##      8.5735310    0.9668330   -0.05607219    0.4640845
##     13.6393255    1.0512265   -0.07560626    0.3615667
##     15.0304471    1.0726835   -0.08764250    0.3353827
##     15.1085159    1.0743573   -0.08879371    0.3342717
##     15.1087870    1.0743641   -0.08879778    0.3342689
##     15.1087870    1.0743641   -0.08879778    0.3342689
##      concave.points_mean symmetry_mean fractal_dimension_mean
## res      0.000000    0.0000000    0.0000000
##      1.015298    0.1221843    -0.04918957
##      1.522512    0.2502993    -0.19391377
##      1.842827    0.3972425    -0.33787796
##      1.952803    0.5017934    -0.43708438
##      1.987360    0.5565752    -0.50011707
##      2.128872    0.6070663    -0.55603567
##      2.201789    0.6322255    -0.58020461
##      2.206636    0.6341719    -0.58195424
##      2.206653    0.6341794    -0.58196081
##      2.206653    0.6341794    -0.58196081

```

We compare the results of using the `glm` function and our `NewtonRaphson` function.

```

tibble(
  predictor = c("(Intercept)", names(Training)[-1]),

```

```

ours = res[nrow(res), -c(1, 2)],
glm = glm(diagnosis ~ ., family = binomial(link = "logit"), data = Training)$coefficients
) %>%
knitr::kable()

```

predictor	ours	glm
(Intercept)	0.4117085	0.4117085
radius_mean	-18.3611026	-18.3611026
texture_mean	1.6021463	1.6021463
perimeter_mean	8.6625595	8.6625595
area_mean	15.1087870	15.1087870
smoothness_mean	1.0743641	1.0743641
compactness_mean	-0.0887978	-0.0887978
concavity_mean	0.3342689	0.3342689
concave.points_mean	2.2066530	2.2066530
symmetry_mean	0.6341794	0.6341794
fractal_dimension_mean	-0.5819608	-0.5819608

Compute the test AUC. (should NOT be used for model comparison)

```

betavec.logit <- res[nrow(res), 3:ncol(res)]
# test data
X_test <- cbind(rep(1, nrow(Test)), model.matrix(diagnosis ~ ., Test)[, -1])
y_test <- Test$diagnosis
# AUC
u_test <- X_test %*% betavec.logit
phat_test <- sigmoid(u_test)[, 1]
roc.logit.test <- roc(response = y_test, predictor = phat_test)
auc.logit.test <- roc.logit.test$auc[1]

```

Resampling on training data: *Does the following resampling method work?*

```
?caret::resamples
```

Hothorn et al. The design and analysis of benchmark experiments. Journal of Computational and Graphical Statistics (2005) vol. 14 (3) pp. 675-699

<https://ro.uow.edu.au/cgi/viewcontent.cgi?article=3494&context=commpapers>

RW-OOB

```

B = 100 # number of bootstrap samples
set.seed(1)
auc.logit <- rep(NA, B)
for (i in 1:B) {
  index_bs <- sample(nrow(Training), replace = TRUE)
  sample <- Training[index_bs, ]
  out <- Training[-index_bs, ]
  res <- NewtonRaphson(dat = sample, func = logisticstuff, start = rep(0, ncol(sample)))
  betavec <- res[nrow(res), 3:ncol(res)]
  X <- cbind(rep(1, nrow(out)), model.matrix(diagnosis ~ ., out)[, -1])

```

```

y <- out$diagnosis
u <- X %*% betavec
phat <- sigmoid(u)[, 1]
roc <- roc(response = y, predictor = phat)
auc <- roc$auc[1]
auc.logit[i] <- auc
}
summary(auc.logit)

```

```

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.9623  0.9764  0.9811  0.9812  0.9863  0.9976

```

```

boxplot(auc.logit)

```

